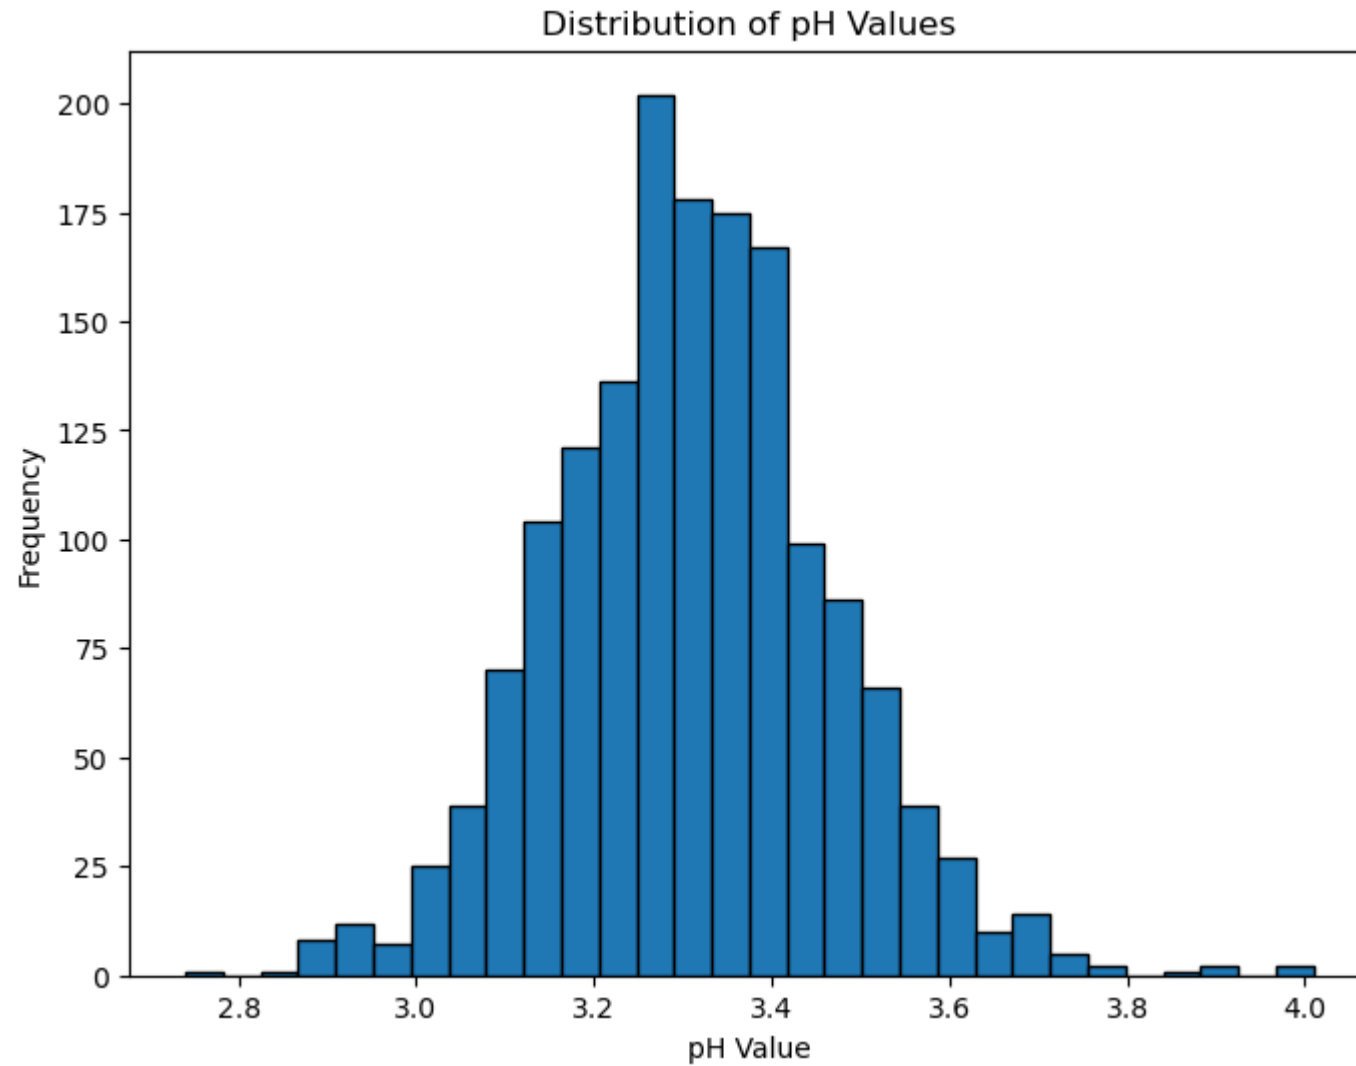
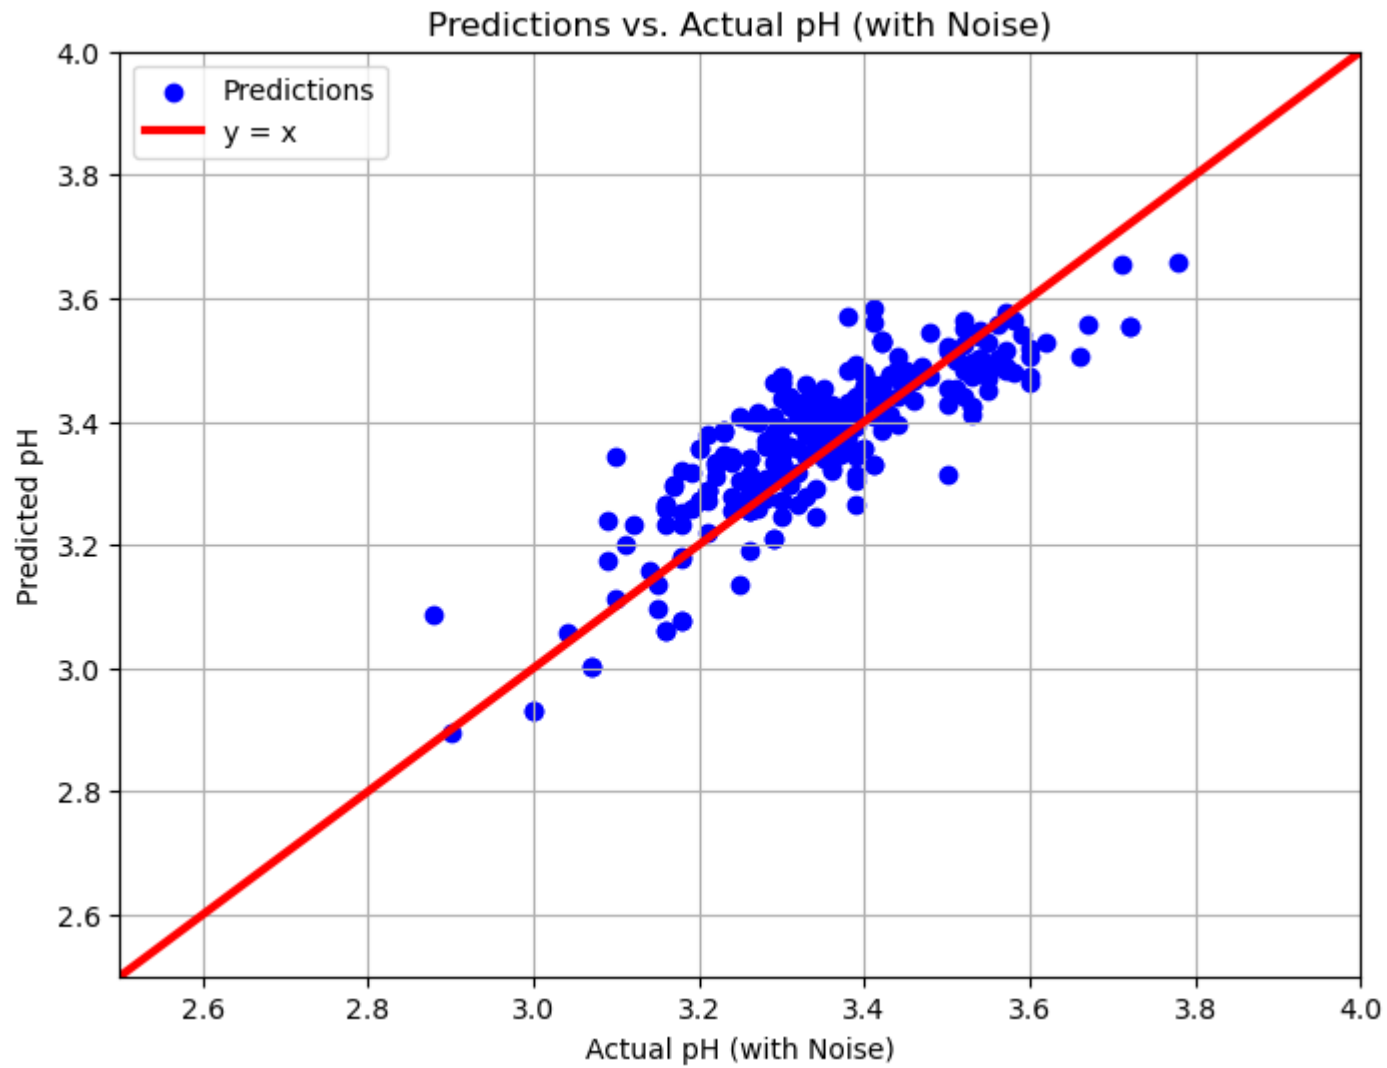





```
In [9]: 1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 # Load the wine quality dataset
7 wine_data = pd.read_csv("winequality-red.csv")
8
9 # Data Preprocessing
10 # Step 1: Handling missing values
11 wine_data = wine_data.dropna()
12
13 # Step 2: Feature selection (choose relevant features)
14 # In this example, we use all columns except 'pH' as features
15 target_column = 'pH'
16 features = wine_data.drop(columns=[target_column])
17
18 # Step 3: Data split
19 # We'll use 80% of the data for training and 20% for testing
20 total_data_points = len(wine_data)
21 train_ratio = 0.8
22 train_data_size = int(total_data_points * train_ratio)
23 test_data_size = total_data_points - train_data_size
24
25 # Split the data into training and test sets
26 X_train = features.iloc[:train_data_size].values
27 y_train = wine_data[target_column].iloc[:train_data_size].values
28 X_test = features.iloc[train_data_size:].values
29 y_test = wine_data[target_column].iloc[train_data_size:].values
30
31
32 # Calculate the weights for the linear regression model using the closed-form solution
33 X = np.concatenate((X_train, np.ones((train_data_size, 1))), axis=1) # Add a column of ones for bias
34 y = y_train
35 X_transpose = X.transpose()
36 weights = np.linalg.inv(X_transpose @ X) @ X_transpose @ y
37
38 # Calculate the mean squared error (MSE) on the test data with noise
39 X_test_with_bias = np.concatenate((X_test, np.ones((test_data_size, 1))), axis=1)
40 predictions = X_test_with_bias @ weights
41 errors = (y_test - predictions) ** 2
```

```
42 mse = np.mean(errors)
43
44 # Visualizations (can be retained for analysis)
45
46 # Visualization 1: Distribution of Wine pH
47 plt.figure(figsize=(8, 6))
48 plt.hist(wine_data['pH'], bins=30, edgecolor='black')
49 plt.xlabel("pH Value")
50 plt.ylabel("Frequency")
51 plt.title("Distribution of pH Values")
52 plt.show()
53
54 # Visualization 2: Predicted vs. Actual pH
55 plt.figure(figsize=(8, 6))
56 plt.scatter(y_test, predictions, c='blue', label='Predictions')
57 plt.plot([2.5, 4.0], [2.5, 4.0], 'r', linewidth=3, label='y = x')
58 plt.xlabel("Actual pH ")
59 plt.ylabel("Predicted pH")
60 plt.legend()
61 plt.grid()
62 plt.title("Predictions vs. Actual pH")
63 plt.xlim([2.5, 4.0])
64 plt.ylim([2.5, 4.0])
65 plt.show()
66
67 # Visualization 3: Correlation Matrix Heatmap
68 correlation_matrix = wine_data.corr()
69 plt.figure(figsize=(10, 6))
70 plt.title("Correlation Matrix Heatmap")
71 sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm")
72 plt.show()
73
74 # Visualization 4: Distribution of Prediction Errors
75 plt.figure(figsize=(8, 6))
76 plt.hist(y_test - predictions, bins=30, edgecolor='black')
77 plt.xlabel("Prediction Error")
78 plt.ylabel("Frequency")
79 plt.title("Distribution of Predictions error")
80 plt.show()
81
82 print("Mean Squared Error :", mse)
83
```

84
85







Mean Squared Error : 0.00710389269362751

In []:

1

